

Default in-training probe was data-starved ($\sim 27\times$ less training data)
Same encoders, same head architecture, only probe-data budget differs

